
Laboratory 3 Report

Process Scheduling Algorithms & Multi-Core Performance Analysis

Course: Operating Systems Lab

Submitted by:

Rhishikesh Tayade : CS24BT041

Himanshu Nagwanshi : CS24BT055

Evaluated Algorithms:

- First-In First-Out (FIFO)
 - Round Robin (RR)
- Multi-Level Feedback Queue (MLFQ)

September 13, 2026

Contents

1	Executive Summary	2
2	Simulator Architecture & Methodology	2
2.1	Core Components and Invariants	2
2.2	Performance Metrics Formulations	3
3	Workload Descriptions	3
4	Experimental Results & Observations	3
4.1	Single-CPU (1 Core) Simulation Results	3
4.2	Dual-CPU (2 Cores) Simulation Results	4
5	Graphical Analysis & Discussion	4
5.1	Average Turnaround Time Comparison	4
5.2	Worst-Case Latency and Maximum Turnaround Time	5
5.3	Round Robin Quantum Sensitivity	5
5.4	MLFQ Priority Boost and Starvation Avoidance	6
5.5	Dual-CPU Multi-Core Acceleration	6
5.6	Execution Timelines (Gantt Schedules)	7
6	Conclusion	7

1 Executive Summary

CPU scheduling forms the computational backbone of multitasking operating systems. This laboratory study investigates the comparative dynamics, trade-offs, and scaling properties of three primary CPU scheduling paradigms:

1. **First-In First-Out (FIFO)**: Non-preemptive, run-to-completion / run-to-block queueing.
2. **Round Robin (RR)**: Preemptive time-slicing across varying quantum sizes ($q \in [1, 100]$).
3. **Multi-Level Feedback Queue (MLFQ)**: 3-tier dynamic priority queue (Q_0, Q_1, Q_2 , each $q = 2$) with priority degradation for CPU hogs, I/O-priority preservation, and periodic priority boosting ($S \in [5, 100]$) to eliminate starvation.

Simulations were conducted under both **Single-Processor (1 CPU)** and **Symmetric Multi-Processor (Dual-Core / 2 CPUs)** configurations across 5 standardized synthetic workload traces.

Key Takeaways:

- **Convoy Effect vs. Interleaving Penalty**: For homogeneous CPU-intensive workloads, FIFO outperforms Round Robin in Average Turnaround Time by up to 26.2% because Round Robin forces all tasks to complete near the end of the total makespan.
- **Quantum Convergence**: As quantum $q \rightarrow \infty$, Round Robin asymptotically converges to FIFO. Fine-grained quanta ($q \leq 2$) increase average turnaround time without benefiting batch jobs.
- **Starvation Mitigation in MLFQ**: MLFQ without boost starves compute-bound jobs in the presence of recurring interactive jobs. A periodic priority boost ($S = 20$) eliminates starvation, restoring deterministic progress.
- **Multi-Core Speedup Limits (Amdahl's Law)**: Dual-CPU systems yield near-ideal speedups of $1.81 \times$ – $1.98 \times$ on compute-bound workloads, but drop to $1.07 \times$ on I/O-bound workloads where hardware devices govern execution latency.

2 Simulator Architecture & Methodology

The scheduling simulator is implemented in C++17 using a tick-by-tick discrete-event execution model advancing time in integer increments ($\Delta t = 1$).

2.1 Core Components and Invariants

1. **Process Model**: Processes maintain static sequences of alternating CPU and I/O bursts terminating with a sentinel -1 . Dynamic state variables track the remaining time in the active burst, current MLFQ queue level, and CPU burst sequence number.
2. **Shared Ready Queue Architecture**: Under Dual-CPU configurations, both cores dequeue processes from a shared ready queue / MLFQ hierarchy, providing automatic work distribution.
3. **Deterministic I/O Resolution**: Processes completing I/O at timestamp t are sorted by $\langle \text{finish_time}, \text{PID} \rangle$ prior to ready queue re-insertion.
4. **Fast-Forward Optimization**: When all CPU cores and ready queues are idle, the simulation clock jumps directly to $\min(\text{next_arrival}, \text{earliest_io_completion})$, preserving strict equivalence while accelerating execution.

2.2 Performance Metrics Formulations

$$\text{Turnaround Time (TAT)}_i = \text{Completion Time (C}_i\text{)} - \text{Arrival Time (A}_i\text{)} \quad (1)$$

$$\overline{\text{TAT}} = \frac{1}{N} \sum_{i=1}^N \text{TAT}_i, \quad \text{Max TAT} = \max_{1 \leq i \leq N} (\text{TAT}_i) \quad (2)$$

$$\text{Speedup Factor } \mathcal{S} = \frac{\overline{\text{TAT}}_{1 \text{ CPU}}}{\overline{\text{TAT}}_{2 \text{ CPUs}}} \quad (3)$$

3 Workload Descriptions

Workload	Tasks	Workload Characteristics	CPU Demand	I/O Demand	Primary
process1.txt	3	Pure compute-bound, staggered arrivals	600	0	Batch Pro
process2.txt	4	Homogeneous equal CPU & I/O bursts	600	40	Balanced I
process3.txt	4	Micro CPU bursts (1) & long I/O (5)	12	40	Highly Int
process4.txt	3	Heterogeneous burst lengths & arrivals	180	85	Mixed Het
process5.txt	6	Extended concurrency version of process4	300	175	Heavy Cor

Table 1: Synthetic workload benchmark characteristics used in the evaluation.

4 Experimental Results & Observations

4.1 Single-CPU (1 Core) Simulation Results

Workload	Metric	FIFO	RR:2	RR:5	RR:10	RR:20	RR:50	MLFQ	MLFQ:20
process1.txt	Avg TAT	363.33	458.67	456.67	453.33	450.00	430.00	461.33	461.33
	Max TAT	590.00	590.00	590.00	590.00	590.00	590.00	590.00	590.00
process2.txt	Avg TAT	525.00	595.00	591.25	585.00	585.00	525.00	595.00	594.50
	Max TAT	600.00	600.00	600.00	600.00	600.00	600.00	600.00	600.00
process3.txt	Avg TAT	14.50	14.50	14.50	14.50	14.50	14.50	14.50	14.50
	Max TAT	16.00	16.00	16.00	16.00	16.00	16.00	16.00	16.00
process4.txt	Avg TAT	123.00	131.33	128.00	132.00	129.00	123.00	132.33	133.00
	Max TAT	163.00	142.00	144.00	157.00	159.00	163.00	144.00	147.00
process5.txt	Avg TAT	189.00	215.67	211.00	200.00	196.33	189.00	216.50	206.67
	Max TAT	225.00	232.00	230.00	225.00	225.00	225.00	235.00	230.00

Table 2: Single-CPU (1 Core) performance metrics across scheduling algorithms.

4.2 Dual-CPU (2 Cores) Simulation Results

Workload	Metric	FIFO	RR:2	RR:5	RR:10	RR:20	RR:50	MLFQ	MLFQ:20	Speedup
process1.txt	Avg TAT	230.00	244.67	245.00	243.33	243.33	246.67	245.33	245.33	1.58x
	Max TAT	390.00	346.00	345.00	350.00	350.00	340.00	344.00	344.00	–
process2.txt	Avg TAT	275.00	301.00	295.00	295.00	295.00	275.00	301.00	302.00	1.91x
	Max TAT	300.00	304.00	300.00	300.00	300.00	300.00	304.00	304.00	–
process3.txt	Avg TAT	13.50	13.50	13.50	13.50	13.50	13.50	13.50	13.50	1.07x
	Max TAT	14.00	14.00	14.00	14.00	14.00	14.00	14.00	14.00	–
process4.txt	Avg TAT	81.33	79.00	81.33	81.33	81.33	81.33	80.33	81.33	1.51x
	Max TAT	93.00	91.00	89.00	93.00	93.00	93.00	92.00	89.00	–
process5.txt	Avg TAT	104.33	108.33	106.33	106.67	105.33	104.33	108.67	108.00	1.81x
	Max TAT	119.00	119.00	113.00	120.00	117.00	119.00	123.00	121.00	–

Table 3: Dual-CPU (2 Cores) performance metrics and FIFO multi-core speedup factors.

5 Graphical Analysis & Discussion

5.1 Average Turnaround Time Comparison




Figure 1: Average Turnaround Time comparison across Single-CPU and Dual-CPU configurations.

Analysis:

- In compute-heavy workloads (**process1**), FIFO achieves the lowest average turnaround time (363.33) because it allows early jobs to finish quickly. Round Robin ($q = 2$) forces concurrent progression, pushing completion times out (458.67).
- In balanced mixed workloads (**process2**), Round Robin with $q < 50$ prolongs average completion (585 – 595) because processes interleave instead of clearing CPU-I/O loops rapidly. Once $q \geq 50$, RR matches FIFO (525.00).

5.2 Worst-Case Latency and Maximum Turnaround Time




Figure 2: Maximum Turnaround Time (tail latency) comparison across algorithms.

Analysis:

- Under FIFO on `process4.txt`, late-arriving interactive processes are starved behind long-running compute jobs, leading to a high Max TAT of 163.00.
- Round Robin ($q = 2$) and MLFQ preempt long jobs, reducing the maximum turnaround time to 142.00 and 144.00 respectively (a **12.9% improvement in tail latency**).

5.3 Round Robin Quantum Sensitivity




Figure 3: Impact of Time Quantum ($q \in [1, 100]$) on Round Robin Average Turnaround Time.

Analysis: As quantum q increases, the number of context preemptions decreases monotonically. When $q \geq \max(\text{burst.length})$, Round Robin mathematically degenerates into FIFO. Very small quanta ($q = 1, 2$) inflict high turnaround penalties on CPU-bound batches.

5.4 MLFQ Priority Boost and Starvation Avoidance




Figure 4: MLFQ Average Turnaround Time vs. Priority Boost Interval S (Dotted line = No Boost baseline).

Analysis: Without priority boost, CPU-bound processes demoted to Q_2 experience starvation when short interactive jobs enter Q_0 . Periodic boosting ($S = 20$) restores balance, reducing Average TAT from 216.50 to 206.67 in `process5.txt`. Setting S too low ($S \leq 5$) undermines the MLFQ priority hierarchy, converting it to standard Round Robin.

5.5 Dual-CPU Multi-Core Acceleration





Figure 5: Dual-Core Speedup Factor (S) across workloads and algorithms.


Analysis:

- **Compute-Bound Speedup:** Compute-intensive workloads achieve $1.81 \times - 1.98 \times$ speedup, approaching the theoretical maximum ($2.0 \times$).
- **I/O Bottleneck Bound:** In `process3.txt`, speedup is constrained to $1.07 \times$ because I/O devices execute independently outside the CPU core pipeline, validating Amdahl's Law.


5.6 Execution Timelines (Gantt Schedules)



(a) FIFO (1 CPU): Process 1 monopolizes the CPU, stalling P2 and P3 until $t = 30$.



(b) Round Robin $q = 10$ (1 CPU): Preemption allows early interleaved progress for P2 and P3.



(c) MLFQ Boost=20 (2 CPUs): Parallel dual-core execution dynamically prioritizes interactive bursts.

Figure 6: Representative Gantt execution schedules for `process4.txt`.

6 Conclusion

1. **FIFO** minimizes average turnaround time for pure batch workloads with known arrivals, but suffers from severe convoy effects under mixed workloads.
2. **Round Robin** guarantees fairness and responsiveness for interactive systems, though its time quantum must be calibrated (10 – 20 time units) to avoid time-slicing overhead.
3. **MLFQ with Periodic Priority Boost** achieves the optimal trade-off for general-purpose OS kernels by dynamically discovering job burst profiles while preventing starvation.
4. **Multi-Core Systems** provide near-linear scalability for CPU-bound tasks, with overall speedup strictly governed by the workload's I/O ratio as formulated by Amdahl's Law.